

AI6122 Text Data Management & Analysis

Topic: Tolerant Retrieval




This lecture

- **Dictionary** data structures
- “Tolerant” retrieval
 - Wild-card queries
 - Spelling correction
 - Soundex




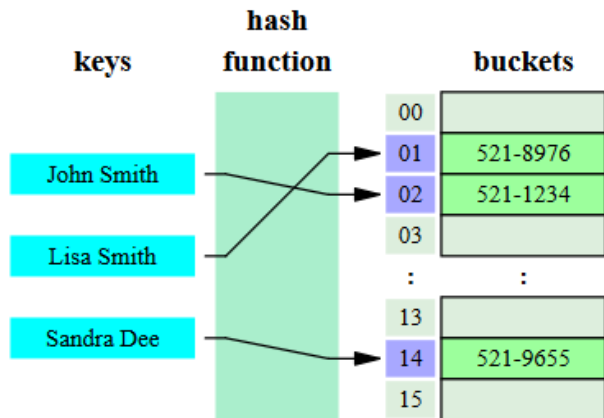
Dictionary data structures for inverted indexes

- The dictionary data structure stores the term vocabulary, document frequency, pointers to each postings list ...
- In what data structure?



Dictionary data structures

- Two main choices:
 - Hash table
 - Tree



Hashes

- Each vocabulary term is hashed to an integer
 - (We assume you've seen hashtables before)
- Pros:
 - Lookup is faster than for a tree: $O(1)$
- Cons:
 - No easy way to find minor variants:
 - judgment/judgement
 - No prefix search [tolerant retrieval]
 - If vocabulary keeps growing, need to occasionally do the expensive operation of rehashing *everything*



Trees

- Simplest: binary tree, or more widely used: B-trees
 - Every internal node has a number of children in the interval $[a, b]$ where a, b are appropriate natural numbers, e.g., $[2, 4]$.
- Trees require a standard ordering of characters and hence strings
- Pros:
 - Solves the prefix problem (terms starting with hyp)
- Cons:
 - Slower than hash
 - Rebalancing binary trees is expensive

“Tolerant” retrieval

- Wild-card queries
 - **mon***: find all docs containing any word beginning “mon”.
- Spelling correction
 - Isolated word
 - Context-sensitive
- Soundex
 - Words with similar pronunciation



Wild-card queries: *

- **mon***: find all docs containing any word beginning “mon”.
 - Easy with binary tree (or B-tree) lexicon: retrieve all words in range: **mon** \leq **w** < **moo**
- ***mon**: find words ending in “mon”: harder
 - Maintain an additional B-tree for terms backwards.
 - Can retrieve all words in range: **nom** \leq **w** < **non**.
- How about ***pro*cent*** ?

Query processing

- At this point, we have an enumeration of all terms in the dictionary that match the wild-card query.
 - We still have to look up the postings for each enumerated term.
- Consider an example query: ***se*ate AND fil*er***
 - This may result in the execution of **many** Boolean *AND* queries.




B-trees handle *'s at the end of a query term

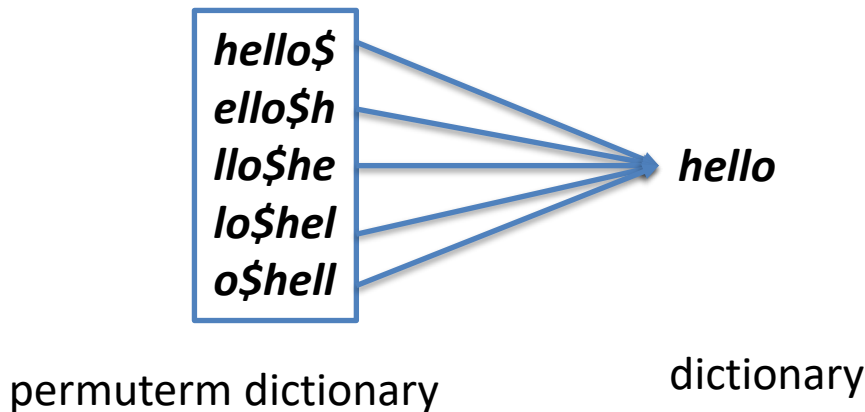
- How can we handle *'s in the middle of query term?
 - *co*tion*
- We could look up **co*** AND ***tion** in B-tree and intersect the two term sets
 - Expensive
- The solution:
 - Transform wild-card queries so that the *'s occur at the end
 - This gives rise to the **Permuterm** Index.

Permuterm index

- Permuterm index is an index of “the terms in the vocabulary”



- For term **hello**, index under: **hello\$, ello\$h, llo\$he, lo\$hel, o\$hell**
 - Symbol \$ is a special symbol, to mark the end of a term



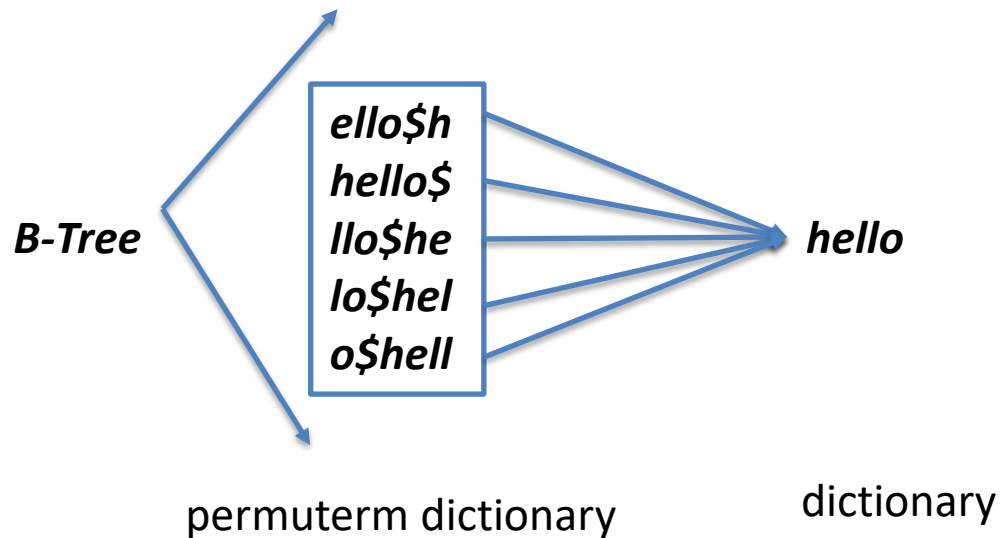
Permuterm index: Examples

- For term **hello**, index under: **hello\$, ello\$h, llo\$he, lo\$hel, o\$hell**

- Queries:

- **X** lookup on **X\$**
 - ***X** lookup on **X\$***
 - **X*Y** lookup on **Y\$X***
- X*** lookup on **X***
 - *X*** lookup on **X***

Query = **hel*o**
X=hel, Y=o
Lookup **o\$hel***



Permuterm query processing


- Rotate query wild-card to the right
- Now use B-tree lookup as before –
 - the terms in the permuterm index are sorted
- Permuterm problem: \approx quadruples lexicon size
 - Empirical observation for English.
- Alternative approach?
 - Bigram (k -gram) indexes

Bigram (k -gram) indexes

- Enumerate all k -grams (sequence of k chars) occurring in any term,
 - From the term “April” we get the 2-grams (*bigrams*)

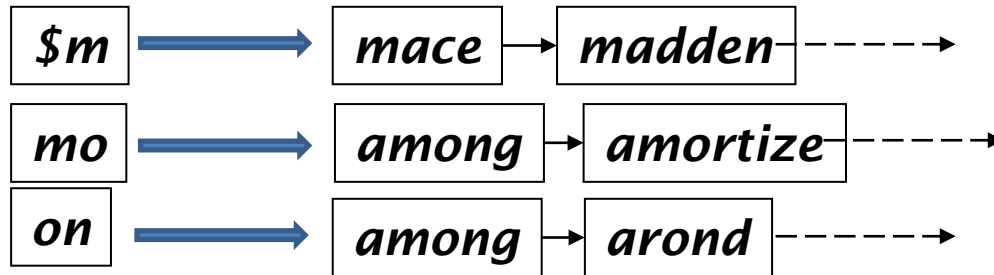
\$a, ap, pr, ri, il, l\$

- \$ is a special word boundary symbol
- Maintain a second inverted index from bigrams to dictionary terms that match each bigram.



Bigram index example

- The k -gram index finds *terms* based on a query consisting of k -grams (here $k=2$).




Processing wild-cards

- Query ***mon**** can now be run as
 - ***\$m AND mo AND on***
- Gets terms that match AND version of our wildcard query.
 - But we'd enumerate ***moon***.
 - Must post-filter these terms against query.
- Surviving enumerated terms are then looked up in the term-document inverted index.
- Fast, space efficient (compared to permuterm).

Processing wild-card queries

- As before, we must execute a Boolean query for each enumerated, filtered term.
- Wild-cards can result in **expensive** query execution
 - very large disjunctions...e.g., `pyth*` AND `prog*`



Spell correction

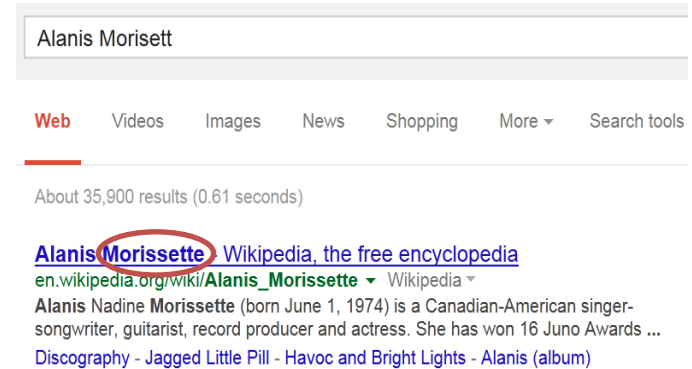
- Two principal uses
 - Correcting document(s) being indexed
 - Correcting user queries to retrieve “right” answers
- Two main flavors:
 - Isolated word
 - Check each word on its own for misspelling
 - Will not catch typos resulting in correctly spelled words
 - e.g., ***from*** → ***form***
 - Context-sensitive
 - Look at surrounding words,
 - e.g., ***I flew form Heathrow to Narita.***

Document correction

- Especially needed for OCR'ed documents
 - Correction algorithms are tuned for this: rn/m
 - Can use domain-specific knowledge
 - E.g., OCR can confuse O and D more often than it would confuse O and I (adjacent on the QWERTY keyboard, so more likely interchanged in typing).
- Goal: the dictionary contains fewer misspellings
- But often we don't change the documents but aim to fix the **query-document mapping**

Query mis-spellings

- Our principal focus here
 - E.g., the query **Alanis Morisett**
- We can either
 - Retrieve documents indexed by the correct spelling, OR
 - Return several suggested alternative queries with the correct spelling
 - Did you mean ... ?



Isolated word correction

- Fundamental premise
 - There is a lexicon from which the correct spellings come
- Two basic choices for this
 - A standard lexicon such as
 - Webster's English Dictionary
 - An “industry-specific” lexicon – hand-maintained
 - The lexicon of the indexed corpus
 - E.g., all words on the web
 - All names, acronyms etc.
 - (Including the mis-spellings)

Isolated word correction

- Given a lexicon and a character sequence Q , return the words in the **lexicon closest** to Q
- What's “**closest**”?
 - Edit distance (Levenshtein distance)
 - Weighted edit distance
 - n -gram overlap



Edit distance and weighted edit distance

- Edit distance: given two strings S_1 and S_2 , the minimum number of operations to convert one to the other
 - Operations are typically character-level
 - Insert, Delete, Substitute
- Weight edit distance: the weight of an operation depends on the character(s) involved
 - Meant to capture OCR or keyboard errors, e.g. **m** more likely to be mis-typed as **n** than as **q**
 - Therefore, replacing **m** by **n** is a smaller edit distance than by **q**
 - This may be formulated as a probability model
- Requires weight matrix as input
- Modify dynamic programming to handle weights

Using edit distances

- Given query,
 - Enumerate all character sequences within a preset (weighted) edit distance (e.g., 2)
 - Intersect this set with list of “correct” words
 - Show terms you found to user as suggestions
- Alternatively,
 - We can look up all possible corrections in our inverted index and return all docs ... slow
 - We can run with a single most likely correction
- The alternatives disempower the user, but save a round of interaction with the user

Edit distance to all dictionary terms?

- Given a (mis-spelled) query – do we compute its edit distance to every dictionary term?
 - Expensive and slow → Alternative?
- How do we cut the set of candidate dictionary terms?
 - One possibility is to use **n-gram overlap**
 - This can also be used by itself for spelling correction.



n-gram overlap

- Enumerate all the *n*-grams in the query string as well as in the lexicon
- Use the *n*-gram index (recall wild-card search) to retrieve all lexicon terms matching any of the query *n*-grams
- Threshold by number of matching *n*-grams
 - Variants – weight by keyboard layout, etc.
- Example:
 - **november** trigram: *nov*, *ove*, *vem*, *emb*, *mbe*, *ber*.
 - **december** trigram: *dec*, *ece*, *cem*, *emb*, *mbe*, *ber*.
 - tri-grams overlap: 3
 - Alternative measure: Jaccard coefficient

Context-sensitive spell correction

- Text: *I flew from Heathrow to Narita.*
- Consider the phrase query “*flew form Heathrow*”
- We’d like to respond

Did you mean “*flew from Heathrow*”?

– because **no docs** matched the query phrase.

Context-sensitive correction

- Need surrounding context to catch this.
- First idea:
 - Retrieve dictionary terms close to each query term (in weighted edit distance)
 - Now try all possible resulting phrases with one word “fixed” at a time
 - *flew from heathrow*
 - *fled form heathrow*
 - *flea form heathrow*
 - Hit-based spelling correction: Suggest the alternative that has **lots of hits**.

Another approach

- Break phrase query into a conjunction of biwords
 - *flew form AND form Heathrow*
- Look for biwords that need only one term corrected.
 - *flew * * form form * * Heathrow*
- Enumerate phrase matches and ... rank them!

General issues in spell correction

- We enumerate multiple alternatives for “Did you mean?”
- Need to figure out which to present to the user
- Use heuristics
 - The alternative hitting most docs
 - **Query log** analysis for especially popular, topical queries
- Spell-correction is **computationally expensive**
 - Avoid running routinely on every query?
 - Run only on queries that matched few docs

Soundex

- Class of heuristics to expand a query into **phonetic** equivalents
 - Language specific – mainly for names
 - E.g., ***chebyshev*** → ***tchebycheff***
- Invented for the U.S. census ... in 1918



Soundex – typical algorithm

- Turn every token to be indexed into a 4-character reduced form
- Do the same with query terms
- Build and search an index on the reduced forms
 - (when the query calls for a soundex match)
- <http://www.creativyst.com/Doc/Articles/SoundEx1/SoundEx1.htm#Top>



Soundex – typical algorithm

- Step 1: Retain the first letter of the word.
- Step2: Change all occurrences of the following letters to '0' (zero):
 - 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'.
- Step 3: Change letters to digits as follows:
 - B, F, P, V \rightarrow 1
 - C, G, J, K, Q, S, X, Z \rightarrow 2
 - D, T \rightarrow 3
 - L \rightarrow 4
 - M, N \rightarrow 5
 - R \rightarrow 6

Soundex continued

- Step4: Remove all pairs of consecutive digits.
- Step5: Remove all zeros from the resulting string.
- Step 6: Pad the resulting string with trailing zeros and return the first four positions, which will be of the form <uppercase letter> <digit> <digit> <digit>.
- E.g., Herman becomes H655.
 - Will *hermann* generate the same code?

Soundex

- Soundex is the classic algorithm, provided by most databases (Oracle, Microsoft, ...)
- How useful is soundex?
 - Not very – for information retrieval
 - Okay for “high recall” tasks (e.g., Interpol), though biased to names of certain nationalities



What queries can we process?

- We have
 - Positional inverted index with skip pointers
 - Wild-card index
 - Spell-correction
 - Soundex
- Queries such as
(SPELL(moriset) /3 toron*to) OR SOUNDEX(chaikofski)