

AI6122 Text Data Management & Analysis

Topic: Event detection



Event detection

- Topic detection and tracking
- Event detection
 - Document-pivot techniques
 - Feature-pivot techniques
- Case study
 - Event detection on Twitter
 - Event detection by queries and documents
 - Event popularity prediction



A bit of history about event detection

- Topic Detection and Tracking (TDT) is a DARPA-sponsored initiative
 - to investigate the state of the art in finding and following new events in a stream of broadcast news stories.
 - TDT Pilot study ran from Sep 1996 to Oct 1997 by DARPA, CMU, Dragon Systems, UMass.
- Three tasks
 - **Segmenting** a stream of data, especially recognized speech, into distinct stories
 - Identifying those news stories that are the first to discuss a new event occurring in the news → **New Event Detection**
 - Given a small number of sample news stories about an event, finding all following stories in the stream → **Event Tracking**



TDT and Event

- TDT: detecting the appearance of new topics and for tracking the reappearance and evolution of them
 - Notion of a “topic” is modified to be an “event” during the study
- **Event:** some unique thing that happens at some point in time.
 - Emphasis more on the “topic” of the event and time, rather than spatial/temporal localization.
 - Example: the eruption of Mount Pinatubo on June 15th, 1991 is considered to be an event, whereas volcanic eruption in general is considered to be a class of events.
 - Events might be unexpected, such as the eruption of a volcano, or expected, such as a political election, or periodical like new year celebration



Where to detect events?

- TDT study assumes **multiple sources of information**, for example various newswires and various news broadcast programs
- The information flowing from each source is assumed to be divided into a **sequence of stories**, which may provide information on one or more events.
 - The general task is to identify the events being discussed in these stories, in terms of the stories that discuss them.
 - Stories that discuss *unexpected events* will of course follow the event,
 - Stories on *expected events* can both precede and follow the event.



TDT Tasks

- *The Segmentation Task*: the task of segmenting a continuous stream of text (including transcribed speech) into its constituent stories.
- *The Detection Task*: **Retrospective Event Detection**
 - The task of identifying **all of the events in a corpus of stories**.
 - Discovering previously unidentified events in an accumulated collection
 - Events are defined by their association with stories.
 - The task is to group the stories in the corpus into clusters. Each cluster represents an event, and the stories in the cluster discuss the event.
 - It will be assumed that each story discusses at most one event. Therefore each story may be included in at most one cluster



TDT Tasks

- *The Detection Task: Online New Event Detection*
 - The task of identifying new events in a stream of stories.
 - Each story is processed in sequence, and a decision is made whether or not a **new event** is discussed in the story after processing the story.
 - The decision is made before processing any subsequent stories (cannot access subsequent stories in online setting)
 - The first story to discuss an event should be flagged YES. If the story doesn't discuss any new events, then it should be flagged NO.
- *The Tracking Task:* associating incoming stories with events known to the system.
 - An event is defined (“known”) by its association with stories that discuss the event. Thus each target event is defined by a list of stories that discuss it.



Event detection

- Event detection nowadays typically refers to the detection of new event and its subsequent stories (i.e., tracking)
 - **Retrospective** Event Detection vs **Online** New Event Detection
- **Document-Pivot Techniques:** event detection is to cluster documents into clusters (events)
 - A document is a data point; event is a cluster
 - Retrospective event detection can use clustering algorithms to access the entire document collection, and to organize the documents into topic clusters, e.g., hierarchical agglomerative clustering (HAC)
 - New event detection: incremental clustering algorithms to process the input streams sequentially,
 - Merge an event with the most similar one,
 - Create a new cluster if the similarity measure exceeds a predefined threshold



Document-Pivot Technique: Incremental clustering

- Take a document d from the document stream (information source)
 - Compute similarity between d and the known events $e \in E$ (i.e., document clusters)
 - If $\text{sim}(d, e) \geq \theta$, assign d to e with the highest similarity
 - If $\text{sim}(d, e) < \theta$, consider d as a new event (a new cluster with a single document for now)
 - Till all documents processed in the stream.
- Parameters to consider:
 - Similarity function, e.g., cosine similarity
 - Document and event representation
 - Tfidf vector? Recent documents in an event be given more weight?
 - Threshold θ
 - Filtering of events: only consider recent events when computing $\text{sim}(d, e)$?



Feature-Pivot Techniques

- Identify topic areas that were previously unseen or rapidly growing in importance within the corpus, **bursty topics**
- Feature-pivot techniques model an event in text streams as a bursty activity, with certain features rising sharply in frequency as the event emerges.
 - An event is therefore conventionally represented by a number of keywords showing burst in appearance counts
 - The underlying assumption is that some related words would show an increased usage as an event occurs.
- These techniques analyze feature distributions and discover events by grouping bursty features with identical trends.



Feature-Pivot Techniques

- “Bursty and hierarchical structure in streams” by Kleinberg (2002)
 - A formal approach for modeling such “bursts”: An infinite-state automaton; Bursts appear as state transitions
 - A nested representation of the set of bursts that imposes a hierarchical structure on the overall stream.
- “Parameter free bursty events detection in text streams” by Fung et al. (2005)
 - Modeled word appearance as binomial distribution, identified the bursty words according to a heuristic-based threshold, and grouped bursty features to find bursty events.
- “Analyzing feature trajectories for event detection” by He, Chang, Lim, and Zhang (2007)
 - Use discrete Fourier transformation (DFT) to categorize features for different event characteristics (e.g., important or not, and periodic or aperiodic events).
 - DFT converts the signals from the time domain into the frequency domain, such that a burst in the time domain corresponds to a spike in the frequency domain

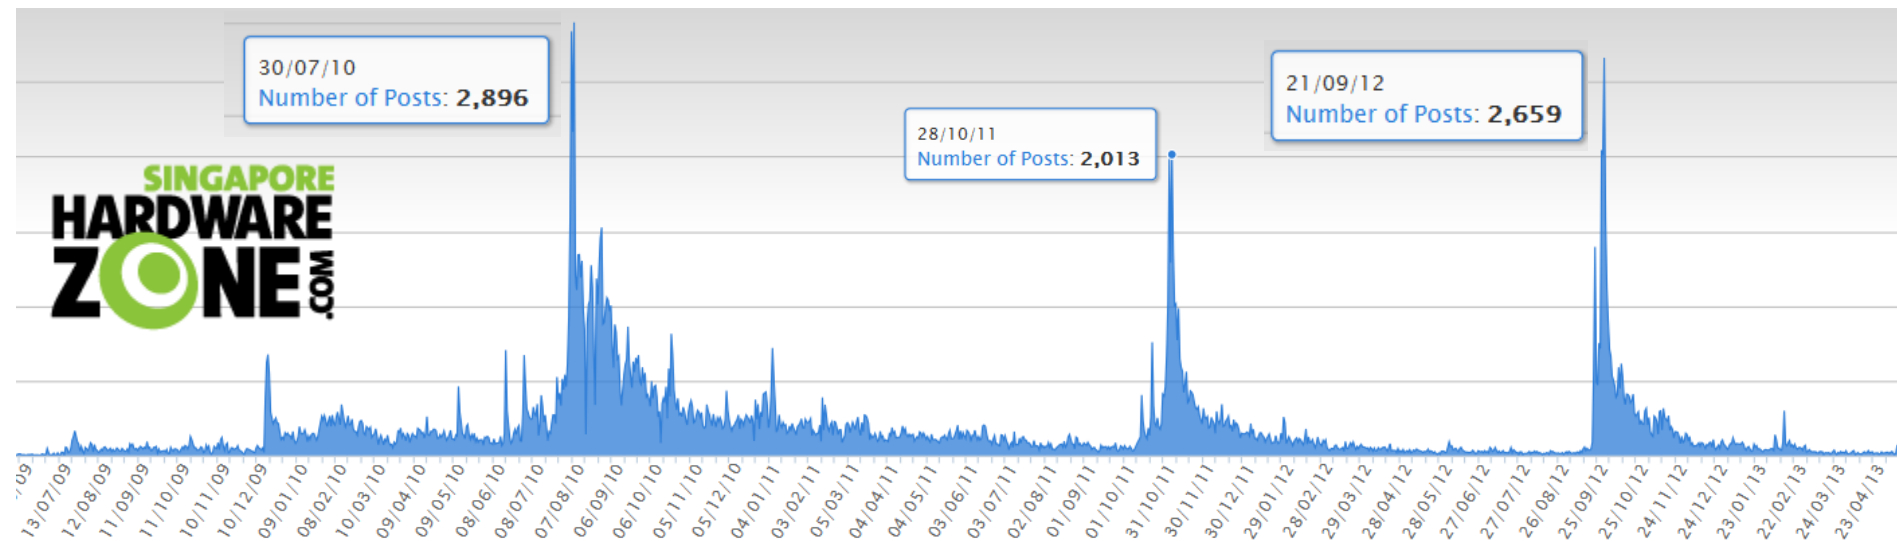


Feature-Pivot Event Detection

- Detect bursty features based on certain models or statistics.
- Events are then detected by maximizing the co-occurrences among documents and the consistence of the frequency distributions for all bursty features within an event.
- The timestamp for an event is calculated based on the bursty periods of the bursty features related to that event



What if a “major event” happens in social media?



Sub-forum: mobile communication technology (2009 – 2013)

Singapore launch date for iPhone 4 is July 30 ... - iMerlion

www.imerlion.com/2010/07/singapore-launch-date-for-iphone-4-is.html ▼

iPhone 4S releases in Singapore on 28 Oct 2011 ...

sgtransport.blogspot.com/.../iphone-4s-releases-in-singapore-on-28.html ▼

SingTel to offer iPhone 5 in Singapore on September 21 ...

info.singtel.com › [About Us](#) › [NewsRoom](#) ▼

13



Event detection applies to Twitter (social media)

- Feature-Pivot Event Detection
 - Identify the bursty words based on certain statistics or model.
 - Grouped bursty words into events based on their co-occurrences.
- Document-Pivot Event Detection
 - Cluster tweets into events. The tweets similar to each other are grouped as events.
 - Certain terms like “named entities” could be assigned with high weights. Named entities need to be recognized by a Named Entity Recognizer
 - Efficient clustering can be achieved through locality sensitive hashing (LSH). Example LSH is MinHash algorithm.
 - The simplest version of the minhash scheme uses k different hash functions, where k is a fixed integer parameter, and represents each set S by the k values of $h_{min}(S)$ for these k functions.



Case study: Segment-based Event Detection from Tweets

- Twitter
 - A message written by the users, up to **140 characters** with **free writing styles**
 - information updates/sharing at low cost
 - A real-time information network that connects you to **the latest information** in your world.
- Event detection in Twitter
 - Events attracted user attentions
 - Events can be more timely detected
- Event detection in Twitter is challenging



Event detection in Twitter: Challenges

- Large data volume
 - 500 million tweets per day in 2019
- Diverse and fast changing topics
- Short and noisy content



PAP POSTERS ARE EVERYWHERE! AND FOR SOME LAMP POLES
THERE ARE BOTH **NSP** AND **PAP** POSTERS!
#whathappentosavingtheearth



ya la! some of them gg to **potong pasir**. I'm gg to **yio chu kang**



Principle of Least Effort [Zipf49]: People used to communicate information with the least context, especially in the situation where a short message with free style is allowed.

- Miss spellings
- Informal abbreviations

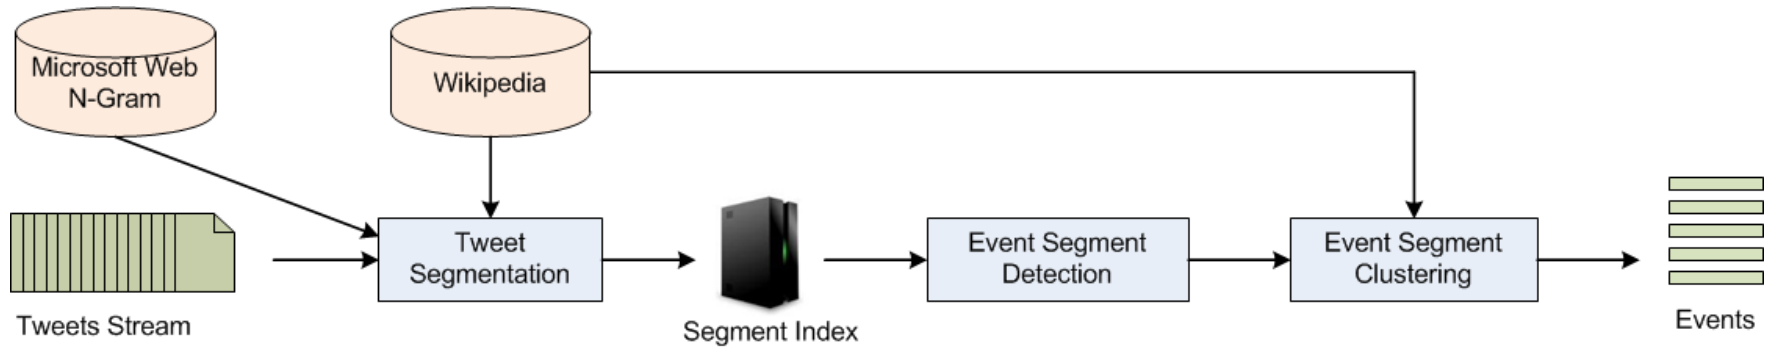


Event detection in Twitter: Our approach

- Tweet Segmentation → **Informative keyphrases**
 - Reduce noise for further processing.
- **User Frequency**
 - Robust to the negative impact of Spam & Self-Promotion tweets.
- External Knowledge Base (**Wikipedia**)
 - Resist to the adverse impact of Pointless Babble tweets.
 - Derive interpretable event descriptions.



TwEvent: System architecture



Tweets → Tweet segments → Event segments → Events



Iphone 4g's coming out on 4th july as according t @zoewasabi hmm. My birthday's on the 10th july. I can use iphone as a present!!! :D

iphone 4g | s | **coming out** | 4th july | according | t | hmm | birthday | 10th july
| use | **iphone** | present | d



iPhone 4G is been officially announced today at WDC
iphone 4g | officially announced | today | wdc

Iphone 4g, iphone,
coming out



Tweet segmentation

- Each segment (unigram/multi-gram) may represent a semantic unit
 - Example segments: **Steve Jobs, MTV Movie Awards**
 - Implemented by maximizing the sum of **stickiness of all segments**
- External resources for calculating the stickiness of a segment.
 - Microsoft Web N-Gram: A prior **probability for each segment** in the index of English web pages
 - Wikipedia: the likelihood that a segment **being an anchor text** in Wikipedia pages.



Tweet segmentation: Example



Example Tweet Portion

youth olympic games sailing competition

Possible segmentation 1

(youth) | (olympic games) | (sailing competition)

Possible segmentation 2

(youth olympic games) | (sailing competition)

Possible segmentation 3

(youth) | (olympic games sailing competition)



Segment burstyness

- **Bursty Segment:** A segment s is a bursty segment in time window t if its tweet frequency $f_{s,t} > E[s|t]$
- **Bursty Probability:** $P_b(s, t) \in (0,1]$ indicates the degree of burstyness of a segment s with frequency $f_{s,t}$.
 - $P_b(s, t) = 1$
 $f_{s,t} \geq E[s|t] + 2\sigma[s|t]$
 - $P_b(s, t) = \text{sigmoid} \left(10 \times \frac{f_{s,t} - (E[S|t] + \sigma[S|t])}{\sigma[S|t]} \right)$
 $f_{s,t} \in (E[s|t], E[s|t] + 2\sigma[s|t])$

Misspelling words and informal abbrev. are detected as bursty segments



Event segment detection

- **User frequency**
 - The number of users who post tweets containing segment s during the time window t .
- Weight each bursty segment: $w_b(s, t) = P_b(s, t)\log(u_{s,t})$
- **Event segment:** A bursty segment s is a potential event segment in time window t if it is ranked among top- K bursty segments by $w_b(s, t)$ where $K = \sqrt{N_t}$
 - N_t is number of tweets in time window t



Event segment similarity

- For each time window t , we further divide the period evenly into M sub-time-window with a weight: $w_t(s, m) = \frac{f_t(s, m)}{\sum_{m'=1}^M f_t(s, m')}$
- A pseudo document $T_t(s, m)$ is built for each segment s at sub-time window m by concatenating all tweets containing s in that window
- The similarity between a pair of segments s_a and s_b is the weighted cosine similarity with tf.idf scheme.

$$\text{sim}_t(s_a, s_b) = \sum_{m=1}^M w_t(s_a, m) w_t(s_b, m) \text{sim}(T_t(s_a, m), T_t(s_b, m))$$

The semantic of a segment is defined by the tweets containing it.



Event segment clustering: k-Nearest Neighbor graph

- Two event segments are in the same cluster if they appear in each others' k-nearest neighbors.
 - An edge between two event segments is retained if and only if they appear in each others' k-nearest neighbors.
 - The resulted connected components are considered as **Candidate Events**.
- **Event**: “anything that happens, especially something important and unusual” --- Cambridge Dictionaries Online
 - Example candidate event: [*Friday night, Friday, weekends, trip, enjoy*] → plans or schedule for weekends
 - **Important and unusual event?**



Event newsworthiness

- **Segment newsworthiness**: the probability that a sub-phrase in the segment appear as anchor text in Wikipedia articles that containing the segment:

$$\mu(s) = \max_{l \in s} e^{Q(l)} - 1$$

- $Q(l)$ is the prior probability that l appears as anchor text in Wikipedia articles that contain l , and l is any sub-phrase of s

- **Event newsworthiness** :

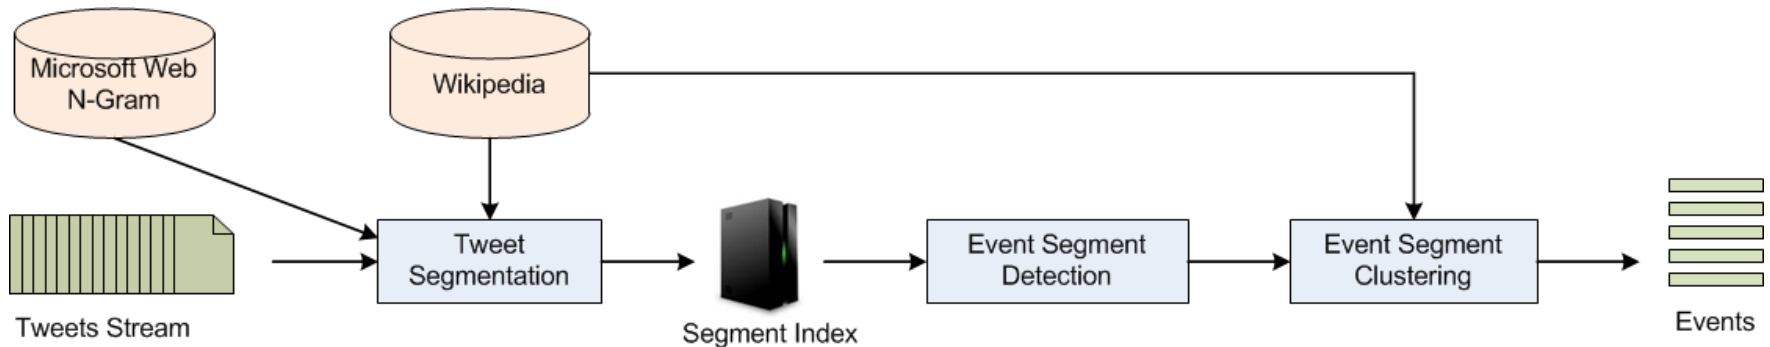
- Often used as anchor text in Wikipedia (well known entities)
- Well connected

$$\mu(e) = \frac{\sum_{s \in e_s} \mu(s)}{|e_s|} \cdot \frac{\sum_{g \in E_e} \text{sim}(g)}{|e_s|}$$



Threshold-based event selection

- Events are selected based on event newsworthiness score from all candidate events in the time window

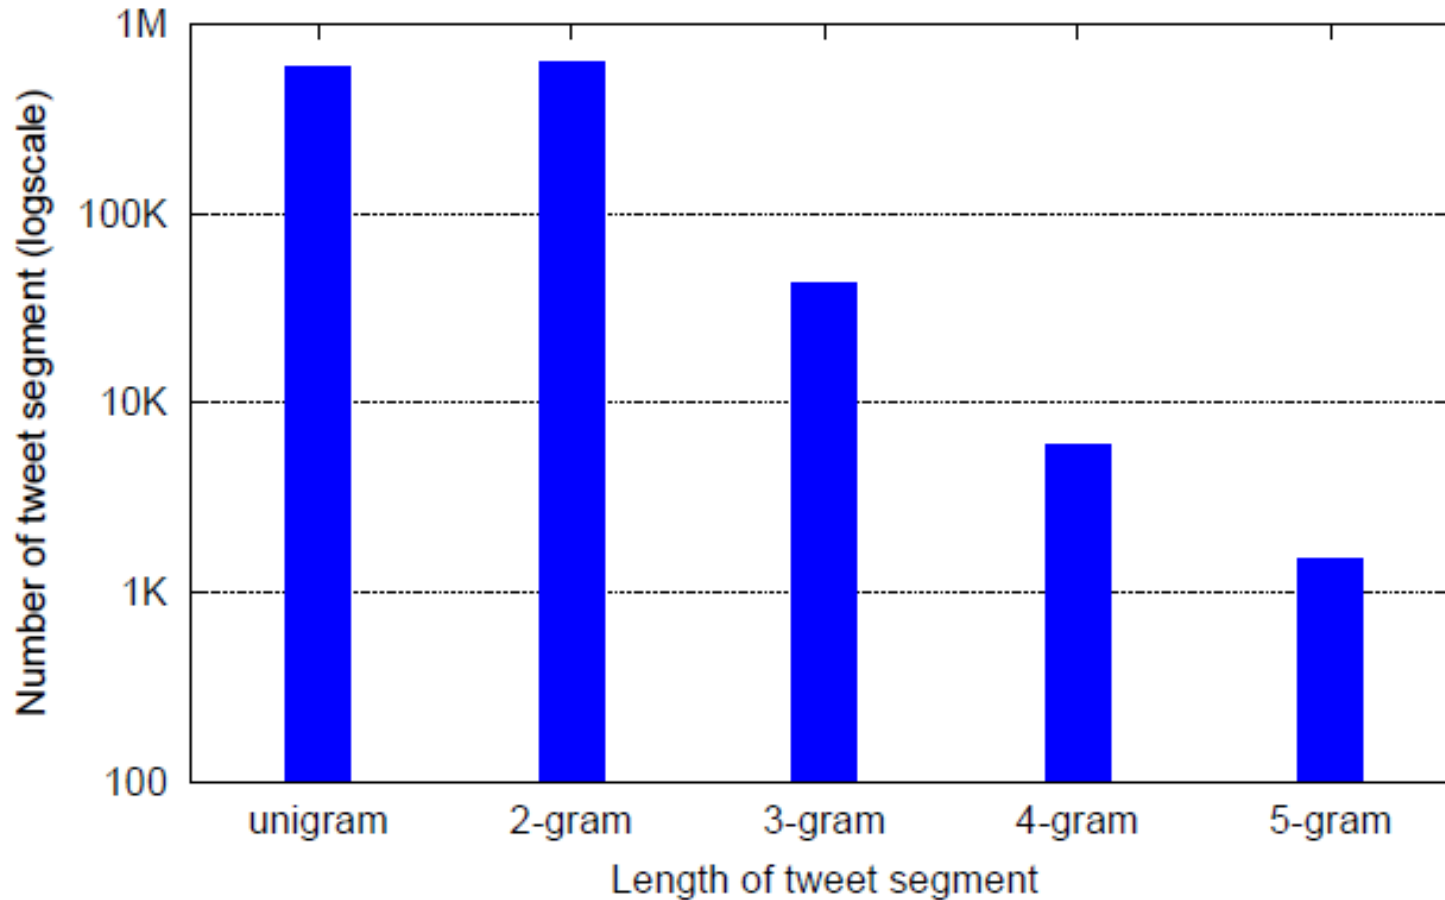


Experiments: Dataset

- Wikipedia:
 - English Wikipedia Dump 2010
- Tweets:
 - 4,331,937 tweets posted in June 2010 by Singapore-based Users.
- Realistic events in data collection period:
 - FIFA World Cup 2010;
 - WWDC 2010;
 - MTV Movie Awards 2010.



Experiments: statistics on segments



Experiments: example events detected

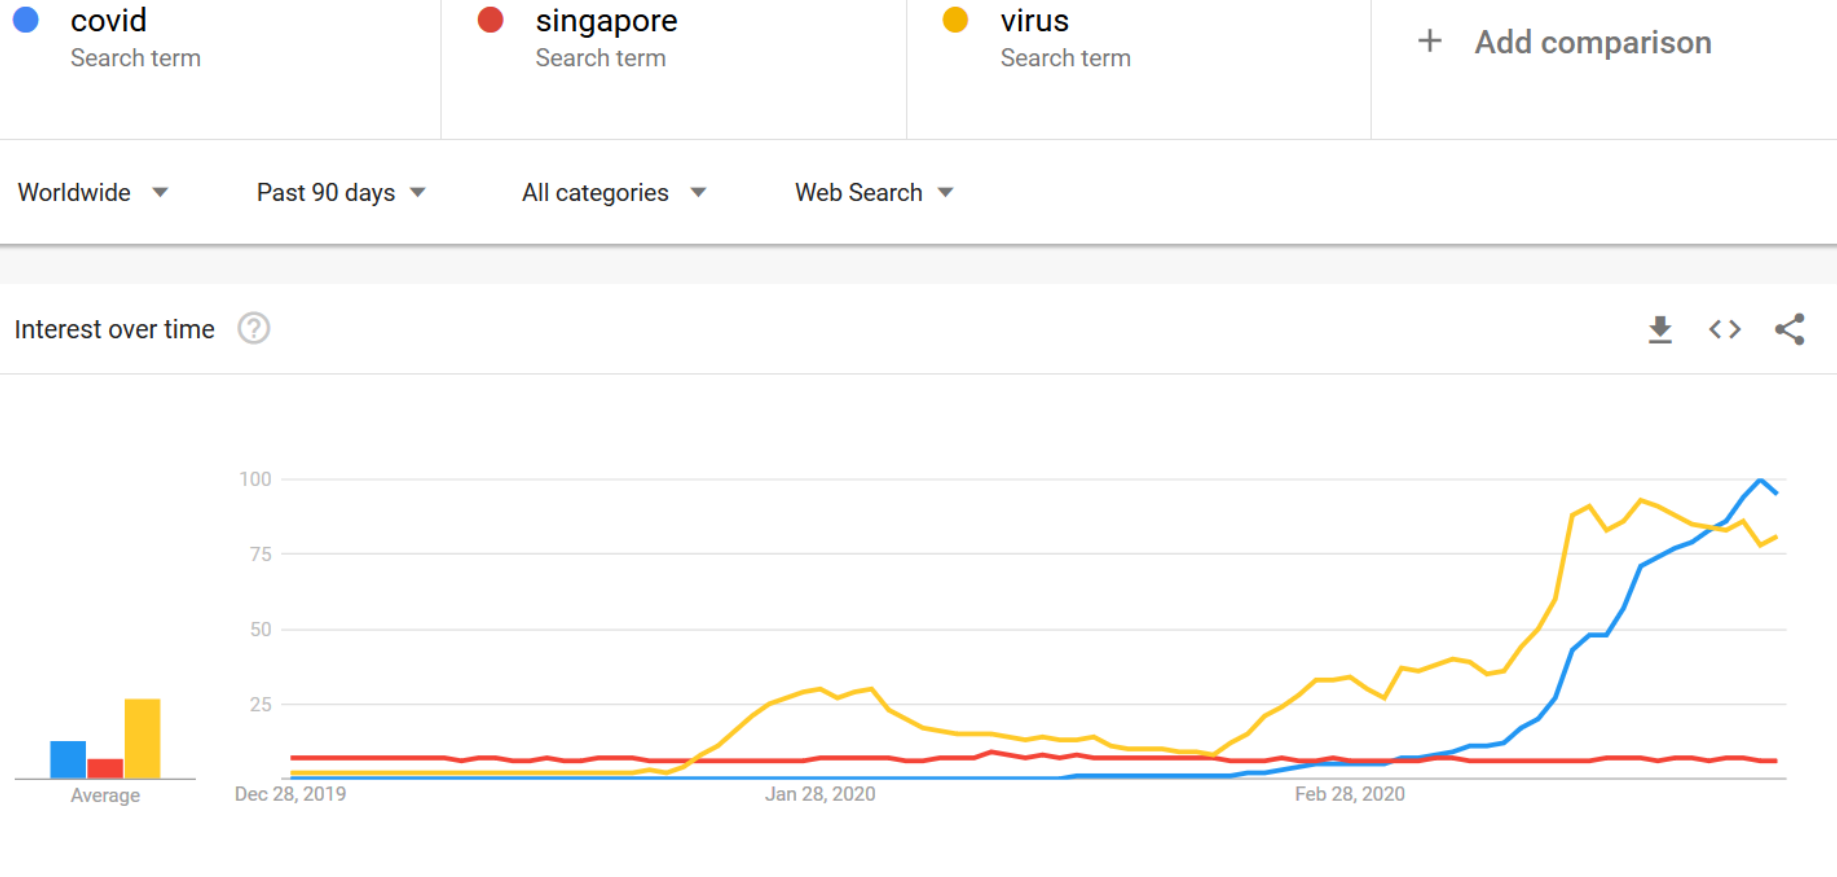
| Day | e_{ID} | [Event Segments]: Event Description |
|-----|------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 7 | e_1 . | [steve jobs, imovie, wwdc, iphone, wifi] : iPhone4 was released during WWDC 2010. |
| | e_2 . | [mtv movie awards, mtv, new moon, twilight, robe] : The movie <i>The Twilight Saga: New Moon</i> was the biggest winner in MTV Movie Awards 2010; it took 4 out of 10 "Best" Awards. |
| | e_3 . | [yesung, yesung oppa, kyuhyun, oppa, kyu] : Korean popular band Super Junior's showcase was held on June 6, 2010 at Singapore. Yesung Oppa and Kyuhyun Oppa are members of Super Junior. |
| 8 | e_4 . | [lady gaga, music video, gaga, mv, alejandro] : The music video <i>Alejandro</i> by Lady GaGa was premiered officially on June 8, 2010. |
| | e_5 . | [ss501, indonesia, ariel, sama, trend] : No clear corresponding real-life event. |
| | e_6 . | [singapore, iphone 4g, iphone 3gs, iphone, coming out] : Related to event e_1 . People started to talk about the release date of iPhone 4 in Singapore. |
| 9 | e_7 . | [lady gaga, youtube, youtube video, music video, gaga] : Related to event e_4 . |
| | e_8 . | [twitter, whale, stupid, capacity, over again] : A number of users complained they could not use twitter due to over-capacity. A logo with whale is usually used to denote over-capacity. |
| | e_9 . | [ipad, iphone, apple, new] : Related to event e_1 . |
| | e_{10} . | [watching glee, glee, season finale, season, channel] : The season finale of the American TV series <i>Glee</i> was broadcasted on June 8, 2010. |
| 10 | e_{11} . | [lady gaga, youtube, youtube video, music video, amber] : Related to event e_7 . |
| | e_{12} . | [justin bieber, try, pa, took, each] : Related to event e_{15} . The song <i>Never Say Never</i> by Justin Bieber serves as the theme song for the movie <i>The Karate Kid</i> , which was released on June 10, 2010 in Singapore. |
| | e_{13} . | [yesung, tweeted] : Super Junior's Yesung posted a photo about his pet turtles. |
| | e_{14} . | [twitter, whale, stupid, capacity, over] : Related to event e_8 . |
| | e_{15} . | [karate kid, watch movie, movie] : The movie <i>The Karate Kid</i> was released on June 10, 2010 in Singapore. |
| 11 | e_{16} . | [uruguay vs france, uruguay, france, vs] : A match between Uruguay and France in World Cup 2010. |
| | e_{17} . | [south africa, vs mexico, mexico, goal, first goal] : A match between South Africa and Mexico in World Cup 2010. And the first goal of the 2010 World Cup was scored in the match. |

Summary on event detection from Twitter

- Tweet Segmentation → **Informative keyphrases**
 - Reduce noise for further processing.
- **User Frequency**
 - Robust to the negative impact of Spam & Self-Promotion tweets.
- External Knowledge Base (**Wikipedia**)
 - Resist to the adverse impact of Pointless Babble tweets.
 - Derive interpretable event descriptions.

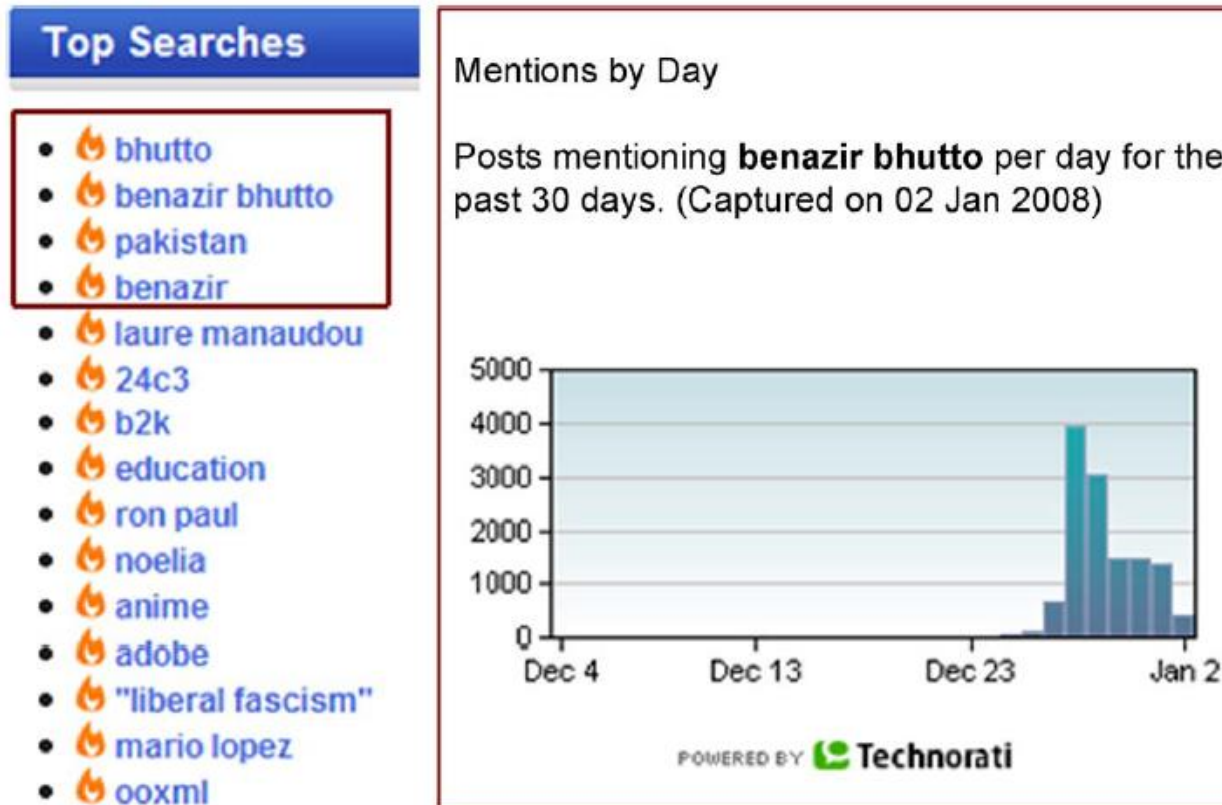


Case study: event detection of common interest



Example

- Top-15 popular searches from Technorati.com captured on December 28, 2007 and statistics on blog posts captured on January 2, 2008



Case study: event detection of common interest

- Events of common interest to many users
 - A large volume of event-related queries are issued to news/blog search engines, making them popular queries during the event period.
 - A large number of news articles and blog posts are published by journalists and bloggers containing updated facts, commentary or discussions about the event.
- From the updated information, web users may formulate new queries (e.g., another person involved in the event) which may subsequently become popular queries.
 - The changes in the queries at different time points become good indications of event evolution.



User in the loop

- Event detection to consider interactions between
 - *query streams*: what people search for
 - *news streams*: what are reported
 - *blog streams*: what are written by users
- Basically: what web users *want to know about* and what they *talk about*.
- Event detection guided by user queries



Challenges

- Not all popular queries issued by masses of web users are event-related.
 - Event-related queries increase dramatically when an event happens does not necessarily imply that all popular queries are event-related.
 - Many extremely popular queries are likely to be website names, such as Google, MySpace, and YouTube, and they are often not event-related.
- Multiple query keywords may be related to the same event.
 - The same query keyword Pakistan issued at different time points may refer to different events happened in that country.
- Computational cost
 - Consider the large number of news articles and blog posts accessible online

Top Searches

- 🔥 bhutto
- 🔥 benazir bhutto
- 🔥 pakistan
- 🔥 benazir
- 🔥 laure manaudou
- 🔥 24c3
- 🔥 b2k
- 🔥 education
- 🔥 ron paul
- 🔥 noelia
- 🔥 anime
- 🔥 adobe
- 🔥 "liberal fascism"
- 🔥 mario lopez
- 🔥 ooxml



Query profile

- **Profile** of a query q at time point t is the set of most recently published documents from news (or blog) stream that match q .
 - Consider to use the most recent 50 matching documents to define the meaning of a query term at the current time (in news search or web search)
- **Recency**: query profile *recency* is the averaged time difference between documents in the query profile to the query's issuing time
- **Clarity**: how indicative the words are in the query profile
 - If a query is event-related, its word features describing the event can often distinguish it from the background.



Query Profile Clarify

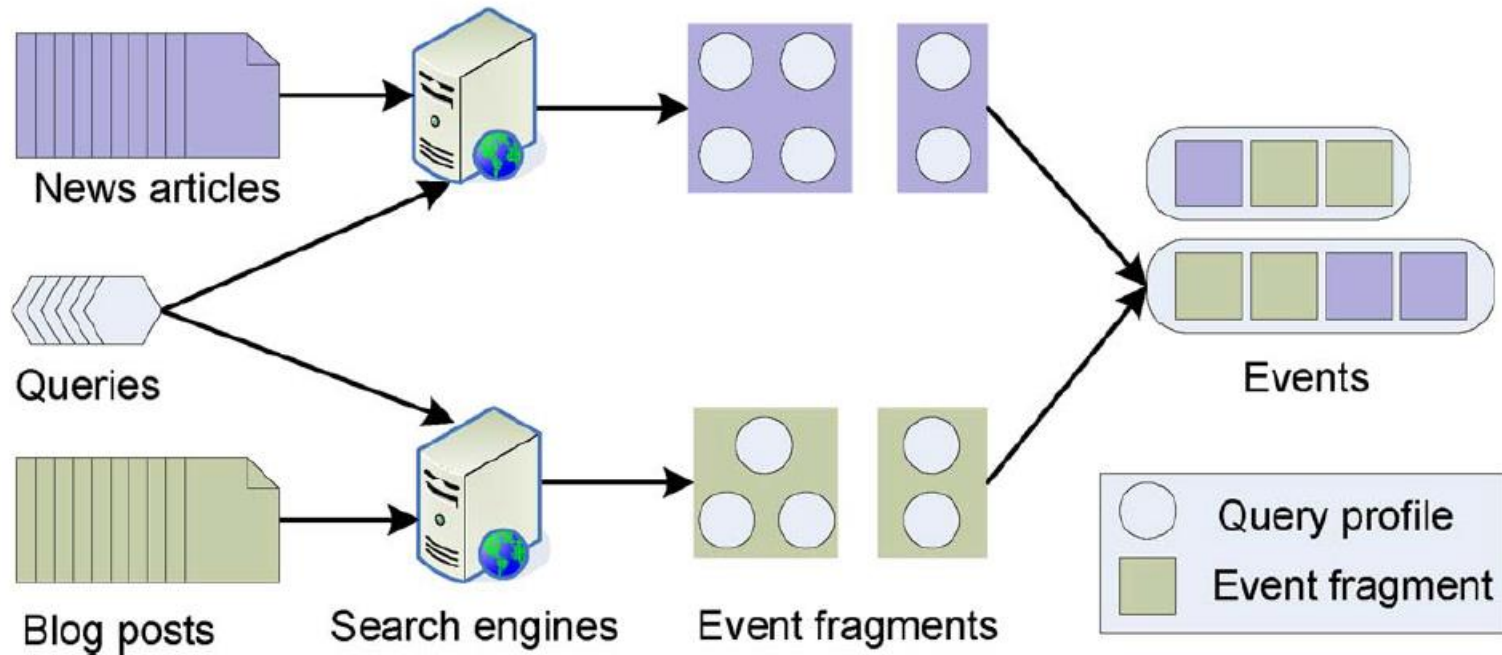
- The word distribution of the query profile, compared to that of a general document collection, measured through Kullback–Leibler (KL) divergence.
 - Divergence between the language model of a query profile, and the language model of a general document collection (i.e., background collection)

$$KL(C_q, S) = \sum_{w \in C_q} P(w|C_q) \log_2 \frac{P(w|C_q)}{P(w|S)}$$

- Assuming query q is about a particular event, then the set of words that are frequently observed among documents in C_q are likely to be describing the event.



Event detection process



Query profile → Event fragment → Event

- For each query in query stream, two query profiles are constructed from the news and blog streams, respectively
 - The event-related query profiles are further passed to event fragment detection.
 - The non-event related query profiles are dropped.
- An event fragment is a set of query profiles that is about the same event received from the same document stream within a predefined time window T .
 - Documents from different streams may demonstrate different properties (e.g., blog posts are noisier than news articles in general),
 - Event fragment detection may require different parameter settings for different document streams.
- Event fragments from both document streams are grouped into events.
 - An event is a sequence of event fragments from both news and blog streams. An event fragment contains query profiles that each contains documents.
 - A detected event contains queries and news articles/blog posts matching them.



Event Fragment Detection

- Given the query profiles in time window T , event fragment detection is to group the query profiles related to the same event into one event fragment (i.e., a small cluster).
- To perform the grouping
 - An appropriate distance metric between any pair of query profiles; and
 - an appropriate clustering algorithm
- Distance between two query profiles
 - Cosine similarity
 - Divergence between the language models of the two (sets of) documents (e.g., square root of Jensen-Shannon divergence or *JSD*)

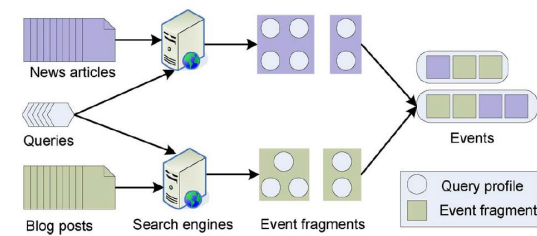


Event fragment detection as clustering

- A clustering algorithm satisfying the following requirements:
 - The algorithm requires no prior knowledge on the number of event fragments to be detected from a set of query profiles,
 - It is unreasonable to guess how many events would happen in a given time window;
 - The algorithm should be able to filter away noise, as predicting whether a query is event-related can never be perfect;
 - The algorithm should be able to handle large data set with reasonable space and time complexity.
- DBSCAN is the choice here



Grouping event fragments into events



- Semantic distance
 - Event fragments in the same event show talk about similar things.
- Query distance
 - Semantic distance worked well for event fragments received from the same stream but not across streams
 - Differences in vocabulary and writing style between news and blogs.
 - Event fragments received within a short time period are likely related to the same event if they share common query keywords.
- Temporal distance
 - An event may last for a long time period and evolve at a fast pace, event fragments of the same event but are temporally far apart may not be similar to each other.
 - Only compute the distance between a newly detected event fragment to those recently detected within 5 days.
 - The timestamp of the event fragment is derived from its query profiles



Case study: Hashtag popularity prediction

- Predict the popularity of new hashtags in the near future (e.g., one day).
 - Popularity range: <25 , $[25, 50)$, $[50, 100)$, $[100, 200)$, >200
- Classification and feature engineering approach
 - Classifiers: NB, kNN, SVM, Logic Regression
 - Features: 7 content features and 11 context features
- Main findings
 - Context features are more effective than content features
 - More effective on bursty tags than continuous tags



7 Content Features and 11 Context Features

| Feature | Description |
|----------------------------------|-------------------------------------------------------------------------------------------------------------------------|
| F_{c1} <i>ContainingDigits</i> | Binary attribute checking whether or not a hashtag contains digits |
| F_{c2} <i>SegWordNum</i> | Number of segment words from a hashtag |
| F_{c3} <i>URLFrac</i> | Fraction of tweets containing URL in T_t^h |
| F_{c4} <i>SentimentVector</i> | 3-dimension vector: ratio of neutral, positive and negative tweets in T_t^h |
| F_{c5} <i>TopicVector</i> | 20-dimension topic distribution vector derived from T_t^h using Topic Model |
| F_{c6} <i>HashtagClarity</i> | KL-divergence of word distribution between T_t^h and tweets collection \mathcal{T} |
| F_{c7} <i>SegWordClarity</i> | KL-divergence of word distribution between tweets containing any segment word in h and tweet collection \mathcal{T} |
| F_{x1} <i>UserCount</i> | Number of users $ U_t^h $ |
| F_{x2} <i>TweetsNum</i> | Number of tweets $ T_t^h $ |
| F_{x3} <i>ReplyFrac</i> | Fraction of tweets containing mention @ |
| F_{x4} <i>RetweetFrac</i> | Fraction of tweets containing RT |
| F_{x5} <i>AveAuthority</i> | Average authority of users in G_t^h |
| F_{x6} <i>TriangleFrac</i> | Fraction of users forming triangles in G_t^h |
| F_{x7} <i>GraphDensity</i> | Density of G_t^h |
| F_{x8} <i>ComponentRatio</i> | Ratio between number of connected components and number of nodes in G_t^h |
| F_{x9} <i>AveEdgeStrength</i> | Average edge weights in G_t^h |
| F_{x10} <i>BorderUserCount</i> | Number of border users |
| F_{x11} <i>ExposureVector</i> | 15-dimension vector of exposure probability $P(k)$ |



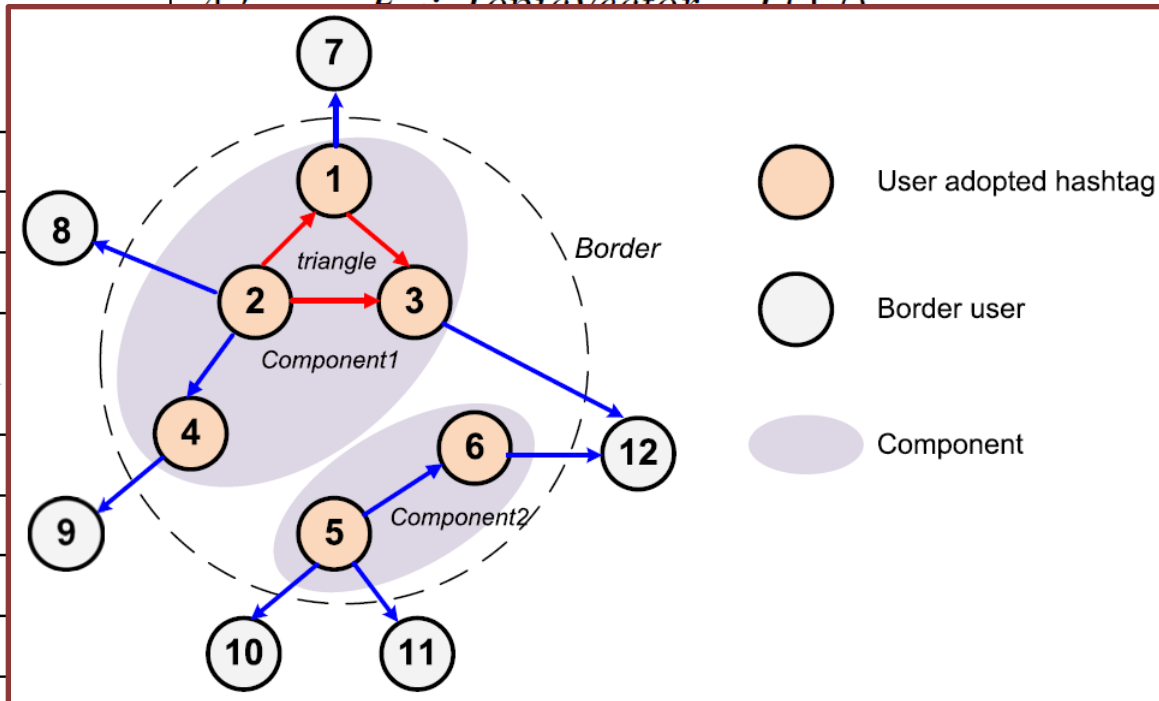
The most and least effective 15 features

| Rank | Feature | Rank | Feature |
|------|------------------------------------|------|--------------------------------|
| 1 | F_{x1} : UserCount | 39 | F_{c5} : TopicVector – T(2) |
| 2 | F_{x10} : BorderUserCount | 40 | F_{c5} : TopicVector – T(14) |
| 3 | F_{x2} : TweetsNum | 41 | F_{x9} : AveEdgeStrength |
| 4 | F_{c6} : HashtagClarity | 42 | F_{c5} : TopicVector – T(17) |
| 5 | F_{x6} : TriangleFrac | 43 | F_{x8} : ComponentRatio |
| 6 | F_{x11} : ExposureVector – P(15) | 44 | F_{c5} : TopicVector – T(20) |
| 7 | F_{x11} : ExposureVector – P(14) | 45 | F_{c5} : TopicVector – T(9) |
| 8 | F_{x11} : ExposureVector – P(9) | 46 | F_{c5} : TopicVector – T(1) |
| 9 | F_{x11} : ExposureVector – P(10) | 47 | F_{c4} : PosRatio |
| 10 | F_{c5} : TopicVector – T(13) | 48 | F_{x5} : AveAuthority |
| 11 | F_{x11} : ExposureVector – P(11) | 49 | F_{c4} : NegRatio |
| 12 | F_{x11} : ExposureVector – P(5) | 50 | F_{c7} : SegWordClarity |
| 13 | F_{x11} : ExposureVector – P(8) | 51 | F_{c4} : NeuRatio |
| 14 | F_{x11} : ExposureVector – P(7) | 52 | F_{c2} : SegWordNum |
| 15 | F_{x11} : ExposureVector – P(12) | 53 | F_{c1} : ContainingDigits |



The most and least effective 15 features

| Rank | Feature | Rank | Feature |
|------|-------------------------------|------|--------------------------------|
| 1 | F_{x1} : UserCount | 39 | F_{c5} : TopicVector – T(2) |
| 2 | F_{x10} : BorderUserCount | 40 | F_{c5} : TopicVector – T(14) |
| 3 | F_{x2} : TweetsNum | 41 | F_{x9} : AveEdgeStrength |
| 4 | F_{c6} : HashtagClarity | 42 | F_{c5} : TopicVector – T(17) |
| 5 | F_{x6} : TriangleFrac | | |
| 6 | F_{x11} : ExposureVector | | |
| 7 | F_{x11} : ExposureVector | | |
| 8 | F_{x11} : ExposureVector | | |
| 9 | F_{x11} : ExposureVector | | |
| 10 | F_{c5} : TopicVector – T(1) | | |
| 11 | F_{x11} : ExposureVector | | |
| 12 | F_{x11} : ExposureVector | | |
| 13 | F_{x11} : ExposureVector | | |
| 14 | F_{x11} : ExposureVector | | |
| 15 | F_{x11} : ExposureVector | | |



Event detection

- Topic detection and tracking
- Event detection
 - Document-pivot techniques
 - Feature-pivot techniques
- Case study
 - Event detection on Twitter
 - Event detection by queries and documents
 - Event popularity prediction



Reference

- Topic Detection and Tracking Pilot Study Final Report
<http://ciir.cs.umass.edu/pubfiles/ir-137.pdf>
- A Survey of Techniques for Event Detection in Twitter.
<https://doi.org/10.1111/coin.12017>
- Case studies:
 - *Twevent: segment-based event detection from tweets*
 - *Query-Guided Event Detection From News and Blog Streams*
 - *On predicting the popularity of newly emerging hashtags in Twitter*

